

Avg pending pods across 10 runs at 400 RPS
(error bar = +/-1 sample stddev across runs)

